

FTD as a Wilsonian Effective Field Theory: Measured β -Function, Ward Identities, Operator Spectrum, and the Limits of Dynamical Standard-Model Emergence

Foundational Ternary Dynamics Research Program
docs/theory/10_eft_program/

2026-04-19

Abstract

We report the first measurement campaign of Wilsonian effective-field-theory observables on the Foundational Ternary Dynamics (FTD) lattice. Five pre-registered experiments, committed to the repository before any code ran, cover: (i) rotational-anisotropy recovery in direction-class correlators; (ii) Lorentz-covariance collapse of temporal against rescaled-spatial flux-flux correlators at $c = 1/\sqrt{3}$; (iii) Ward-identity closure for the Gauss constraint and the composite correlator $\langle(\nabla \cdot J)J^\nu\rangle$; (iv) a three-scale coupling extraction yielding a measured β -function from real-space block-spin blocking; and (v) a scaling-dimension measurement for six gauge-invariant operators. We additionally run three dynamical Standard-Model emergence tests (electroweak symmetry breaking cold-start, three-generation cold-start, and a large-L extrapolation scan of α_{eff}).

The measurements yield partial successes, documented failures, and one result that subsequent analysis (§11) reframes from “QED-running” to “lattice Coulomb geometry,” all reported without post-hoc adjustment of the pre-registered expectations. Lorentz-covariance residuals reach 0.4% at $r \leq 4$ lattice units and grow to 5% at $r = L/4$, consistent with the cubic-lattice dispersion correction $\omega = 2c \sin(k/2) - ck \sim \mathcal{O}(k^2)$. Ward-identity closure is SOR-tolerance-limited at $\sim 1\%$ of $|J|_{\text{max}}$ in the production solver, later improved to $\leq 10^{-8}$ on deep vacuum by a matched-stencil CG solver (§10). The measured “ $\alpha_{\text{eff}}(L)$ ” plateau — originally framed as a β -function running — is re-identified in §11 as the periodic lattice Poisson Green’s function $\alpha_r(r, L) = 2rG_L(r)$ at $r/L \approx 0.31$, a zero-free-parameter prediction whose match to engine measurement is $R^2 = 1.0000$ in the Coulomb tail. The original “ $160\times$ QED β ” claim is retracted: that channel carries no fine-structure content. Six operators give high- R^2 scaling-dimension fits but all cluster near $\Delta \approx 0.5$ rather than the pre-registered naive-counting values [2, 5], attributed to a pulse-envelope artefact in the measurement scenario. Cold-start electroweak symmetry breaking is not observed at the lattice sizes tested ($L \leq 16$); the Higgs vacuum expectation value remains [IMPOSED] in the parametric-insertions catalogue.

The campaign ends with FTD characterised as *an emerging Wilsonian EFT*: the five-pillar checklist (Ward, Lorentz, RG, operator expansion, large-L matching) produces measurable signals in the expected qualitative directions, but quantitative agreement at the largest tested L demands larger lattices and higher-tolerance gauge projection than the present engine implements. We enumerate the concrete follow-up tickets that would close each gap.

A Day-2 campaign (§10) closes four of those tickets: matched-stencil CG Poisson drops the Ward-identity floor from $\sim 10^{-2}$ to $\leq 10^{-8}$ on deep vacuum (a million-fold improvement); an EWSB amplitude threshold map locates a first-order condensation transition in $(0.6, 0.7)$ with spontaneous charge-sign selection above it; spectroscopy on the condensate gives clean mass gaps ($m \approx 0.18$ at $\text{amp} = 0.80$) with two independent channels agreeing to 3%; and Rutherford scattering cross-checks the $V(r)$ extraction at small impact parameter, returning

the same lattice-Coulomb-geometry plateau (now interpreted per §11) consistent with the static measurement.

A Phase-F follow-up (§11) unblocks the GPU build via WSL2 and runs a 4-point r-max sweep at $L = 64, 128, 256, 384$ under consistent conditions (ticks = 150, same $r/L \approx 0.31$ slice), and a line-by-line audit of the extraction pipeline (`AUDIT_ALPHA_EXTRACTION.md`) followed by a Phase-G analytical resolution (`DERIV_EMERGENT_COULOMB_GEOMETRIC.md`) establishes the honest interpretation. The measured $\alpha_r(r, L)$ is *not a coupling constant at all*: the engine’s emergent-mode Gauss law is $\nabla \cdot J = s$ with no coupling constant, so $\alpha_r(r, L) = 2rG_L(r)$ is a parameter-free lattice-Green’s-function prediction whose zero-free-parameter match to measurement is $R^2 = 1.0000$ (median 0.07% relative error) across 16 points at $L=384$ in the Coulomb-tail regime. The “plateau at $3.6 \times \alpha_{\text{ref}}$ ” is the value of the periodic lattice Poisson kernel at $r/L \approx 0.31$ on a cubic torus — pure geometry, no fine-structure content. This code path carries no α . A parametric coupling added to Gauss’s law (Phase H, spec’d but not yet measured) would be the appropriate test of whether FTD’s dynamics can emergently reproduce α_{ref} .

1 Introduction

Foundational Ternary Dynamics (FTD) is a discrete lattice framework whose phenomenological outputs have been benchmarked against the Standard Model for three years [1, 2]. The framework’s algebraic spine [4] contains nine theorems including the master quadratic $x^2 - 16G_*^2x + 16G_*^3 = 0$ with $G_* = \Gamma(1/4)/\Gamma(3/4)$ (THEOREM); the polynomial’s larger root $x_+ \approx 137.036$ matches CODATA α^{-1} to 1.26 ppm, with this identification carrying tag [STRONGLY MOTIVATED CONJECTURE] (not [THEOREM]) as supported by structural-uniqueness scans [5, 6] but lacking a derivation chain from axioms to the physical identification. The smaller root $x_- \approx 3.024$ is numerically close to $N_c = 3$ (0.80%); the identification $x_- = N_c$ has since been retired from the canonical ledger (FTD-0014 removed) — $N_c = 3$ is established by independent topological derivations. The Moore-layer decomposition yielding three fermion generations [8] is at [SELECTION]. This paper does not address those algebraic / number-theoretic claims; its scope is the engine’s *dynamical* output and its standing as a Wilsonian effective field theory under the five-pillar checklist: a measured β -function, verified Ward identities, quantified Lorentz-covariance recovery, a classified operator basis, and a controlled large- L extrapolation.

This paper reports that measurement campaign. The experimental programme is pre-registered in `SPEC_EFT_RECOVERY_PROGRAM.md`, committed to the project repository before any new engine code ran. Section 2 fixes the lattice, toggles, and reference regime. Section 3 reports the symmetry-recovery measurements. Section 4 the β -function extraction. Section 5 the operator spectrum. Section 6 the three dynamical SM tests. Section 7 catalogues the updated [IMPOSED] vs [MEASURED] tags. Section 8 enumerates the concrete next steps.

Epistemic discipline. Three rules, stated in SPEC §1, govern every result in this paper: (i) no numerical fishing — every measurement has a pre-registered expectation; (ii) no relabelled insertions — an imported formula with FTD-supplied parameters stays [PARAMETRIC], not [DERIVED]; (iii) every module is ≤ 500 lines with one clear responsibility, every observable has a CTest.

2 Lattice Setup and Canonical Regime

The FTD lattice is a three-dimensional cubic lattice of size L^3 with periodic boundaries. Each voxel carries a ternary state $s \in \{-1, 0, +1\}$ and a continuous flux vector $\mathbf{J} \in \mathbb{R}^3$. The update rule proceeds through five phases per tick: `phase_read` (Laplacian propagation), `phase_write`

(state evolution with optional genesis), `gauss_project` (six SOR iterations enforcing $\nabla \cdot \mathbf{J} = s$), `phase_forces` (Coulomb, gravity, etc. when enabled), `phase_movement` (particle advection).

The CFL stability condition on the cubic lattice gives the maximum information-propagation speed $c = 1/\sqrt{D} = 1/\sqrt{3}$ [1]. The bare coupling is $G_C = \sqrt{\alpha}$ with $\alpha = 1/X_+$ from the master-quadratic root $X_+ = 137.036 \dots$ [3].

Canonical reference regime. Per SPEC §3: $L = 64$, scenario `flux-pulse`, measurement tick $t = 2000$, seed 42, damping OFF during measurement, `gauss_projection` ON. Exceptions are declared at each measurement and documented without retrofit.

3 Symmetry Recovery (Phase 1)

Three measurements test whether the continuous symmetries a Wilsonian EFT requires are emergently respected by the discrete engine.

3.1 Rotational Anisotropy

We compute the flux-flux correlator $C(r) = \langle \mathbf{J}(x) \cdot \mathbf{J}(x+r) \rangle$ separately along the three O_h orbits of lattice displacement vectors: face $(1, 0, 0)$, edge $(1, 1, 0)$, diagonal $(1, 1, 1)$. The anisotropy coefficient $\delta(r) = (\xi_{\text{face}} - \xi_{\text{diag}})/\bar{\xi}$ is extracted from exponential fits to each orbit’s correlator.

Finding. *The fit infrastructure (implemented in `eft/anisotropy.h`) is validated on four configurations: uniform flux (residual $< 10^{-9}$), plane wave (strong expected signal), isotropic Gaussian noise ($\delta/C(0) < 5 \times 10^{-3}$), and synthetic exponential (recovers $\xi = 5.0$ to 0.02% with $R^2 = 0.99998$).*

A realistic-dynamics anisotropy campaign, measuring $\delta(r)$ at $L \in \{32, 48, 64, 96\}$ on the canonical regime, is deferred to a follow-up (Section 8).

3.2 Lorentz Recovery

Temporal and spatial flux-flux correlators are measured on a seeded plane-wave $J_x(z) = \cos(k_z z)$ at $k_z = 2\pi/L$, $L = 64$, over $T = 512$ ticks with damping OFF. After rescaling $\tau \rightarrow r/c$ with $c = 1/\sqrt{3}$ and normalising each correlator to its $C(0)$ value, the two curves should collapse.

Finding. *Residuals $|C_t(r/c) - C_s(r)|/|C_s(r)|$ are 0.4% at $r = 2$, 0.8% at $r = 4$, 1.9% at $r = 8$, rising to 4.7% at $r = 12$ and peaking near the $r = L/4 = 16$ zero-crossing. The pre-registered 1% target is met for $r \leq 4$; the growth to 5% at mid- r is dominated by lattice-dispersion corrections $\omega/k - c \sim -(k/2)^2/6 \approx 1.3 \times 10^{-3}$ plus time-integration error.*

The $\omega - ck$ dispersion residual alone cannot explain the observed magnitude; additional error from the 6-iteration SOR projection and the standing-wave initial condition contribute. Full recovery to permille precision requires $L \geq 128$.

3.3 Ward Identities

Four pre-registered tests on configurations of increasing complexity (vacuum, static charge pair, oscillating dipole, plane-wave EM field) measure $\max_x |\nabla \cdot \mathbf{J}(x) - \rho(x)|$ restricted to vacuum voxels ($s = 0$; the `gauss_project` routine is designed to leave particle voxels unmodified).

Finding. *Vacuum Gauss constraint holds to $< 10^{-9}$ on an empty lattice. With a $+1/-1$ pair at separation 4 and 20 ticks of `gauss_project`, the maximum vacuum residual is 2.6×10^{-2} , with*

RMS 2.8×10^{-3} (i.e. $\sim 10\%$ of $|\mathbf{J}|_{\max}$). The pre-registered machine-precision expectation ($\leq 10^{-8}$) is off by six orders of magnitude.

Investigation (§4 of `DERIV_SYMMETRY_RECOVERY.md`) traces the gap to the fixed 6-iteration SOR implementation in `engine/src/poisson_solvers.cpp`: with $\omega = 1.75$, each sweep reduces the Poisson residual by $\sim 0.5\times$, so six sweeps achieve $\sim 0.5^6 \approx 1.5\%$ — in agreement with measurement. The pre-registration was too optimistic about the engine’s tuning, not about the physics. A one-shot `gauss_project_converged()` would close the gap; the ticket is noted in §8.

The composite Ward identity $\langle (\nabla \cdot \mathbf{J})(x) J^\nu(x+r) \rangle = \langle \rho(x) J^\nu(x+r) \rangle$ is satisfied with RMS residual $\sim 4 \times 10^{-4}$ across the canonical pair configuration — within 1000% of the pre-registered 0.1% threshold.

Vertex Ward identity. $\Gamma_\mu(p, p) = \partial \Sigma / \partial p^\mu$ is documented as [OPEN]. The engine does not yet carry lattice fermion propagators from which self-energy Σ could be measured; this is a post-campaign extension.

4 Measured β -Function from Real-Space Blocking (Phase 2)

We extract the scale-dependent coupling $\alpha_{\text{eff}}(L)$ at three lattice sizes via the two-charge interaction potential $V(r)$. For pure Coulomb $V = -\alpha/r + C$, linear regression of V against $1/r$ gives α directly; on a screened Coulomb $V = -\alpha e^{-r/\lambda}/r$ a nonlinear fit separates coupling from screening length.

4.1 Three-Method Extraction at $L \in \{16, 32, 64\}$

Table 1: α_{eff} extractions across three lattice sizes. All values dimensionless. Ratio to $\alpha_{\text{ref}} = 1/137.036 \approx 0.00730$.

L	slope fit (V vs $1/r$)	Yukawa (α, λ)	tail α_r plateau
64	0.120 ($R^2 = 0.91$)	0.176, $\lambda = 10.6$	0.033 ($r \in [14, 20]$)
32	-0.127 ($R^2 = 0.45$)	0.706, $\lambda = 2.9$	0.035 ($r \in [8, 10]$)
16	invalid (1 r -point only)	invalid	invalid

4.2 $\beta(g)$

With $g = \sqrt{\alpha}$ and $\beta = [g(\text{scale} \cdot 2) - g(\text{scale})]/\ln 2$:

- Tail-plateau method: $\beta_{\text{measured}} = -8.3 \times 10^{-3}$ at $g = 0.182$; ratio to $\beta_{\text{QED}, 1\text{-loop}}(g) = g^3/(12\pi^2) = 5.1 \times 10^{-5}$ is -163 .
- Yukawa method: $\beta_{\text{measured}} = -0.61$; ratio -978 (dominated by method-sensitivity to the screening-length extraction).
- Slope method: insufficient data.

Finding. *[retracted; see §11] β is negative, matching the asymptotic-freedom direction of textbook QED, but its magnitude is roughly 160× the one-loop QED prediction. This is a new quantitative*

~~prediction: the FTD lattice at $L=64$ exhibits β that runs $\sim 160\times$ faster than textbook QED, testable by repeating the measurement on a lattice where $r \gg \lambda_{\text{Yukawa}}$ (i.e., $L \geq 128$).~~

Retraction note. The Phase-F follow-up audit (`AUDIT_ALPHA_EXTRACTION.md`) and the Phase-G analytical resolution (§11, `DERIV_EMERGENT_COULOMB_GEOMETRIC.md`) re-identify this “ β -running” as the periodic lattice Poisson Green’s function $\alpha_r(r, L) = 2rG_L(r)$, a zero-free-parameter geometric quantity with no fine-structure content and no QED-comparable running. The “ $160\times$ QED” framing is a category error, not a quantitative prediction. The native source/flux $\beta_\mathcal{E}$ measured on the Langevin + genesis ensemble (FTD-0070, Phase 2 closure of the bridge contract) is a separate, native-EFT observable, consistent with a Gaussian fixed point at $b \in \{1, 2, 4, 8\}$ within 1σ .

4.3 Charge Conservation Under Blocking

A pre-block validation gate (SPEC §5.1) verifies that the block-spin transformation (factor-of-2 real-space coarse-graining) preserves total charge exactly across 12 test configurations (single positive, charge pair, three distant charges, uniform flux, Gaussian noise, iterated blocking, and an overflow case where a block contains 4 same-sign children). All 12 pass; the charge-conserving variant of `block.state` is validated at machine precision.

5 Operator Spectrum (Phase 3)

Six gauge-invariant operators are measured against their pre-registered naive-dimension brackets:

Table 2: Pre-registered vs measured scaling dimensions. “Bracket” from SPEC_OPERATOR_BASIS.md §3.

Operator	Naive Δ	Bracket	Measured Δ	R^2
$\mathbf{J} \cdot \mathbf{J}$	2	relevant	0.531	0.997
$(\nabla \cdot \mathbf{J})^2$	4	marginal	0.458	0.918
$(\nabla \times \mathbf{J}) \cdot (\nabla \times \mathbf{J})$	4	marginal	0.391	0.959
$\mathbf{J} \cdot \nabla(\nabla \cdot \mathbf{J})$	5	irrelevant	0.563	0.370
$(\mathbf{J} \cdot \mathbf{J})^2$	4	borderline	0.753	0.992
$s \cdot s$	2	relevant	—	(no manifested charges)

Finding. Five of six operators give valid fits with high R^2 ; the measured Δ values cluster near 0.5 regardless of the operator’s naive dimension. This is attributed to a pulse-envelope artefact: the measurement scenario (localised Gaussian flux pulse) imposes a spatial envelope of width $\sim 20a$ that dominates the correlator decay at $r \leq L/4$, obscuring operator-specific scaling.

The Phase-3 contribution is *measurement infrastructure*, not validated operator classifications. Larger lattices or confinement-era scenarios (where long-range structure exists in a controlled background) are needed to test the naive-dim brackets properly.

6 Dynamical Standard-Model Emergence (Phase 4)

Three pre-registered tests of the engine’s ability to produce Standard-Model physics from cold-start initial conditions.

6.1 EWSB Cold-Start

A bare flux seed with uniform amplitude 0.15, genesis ON, evolved for 2000 ticks at $L = 16$. Observable: $\langle |J| \rangle$ and total charge over time.

Finding. $\langle |J| \rangle$ decays to $0.48 \times$ initial amplitude; no charges ever manifest. Branch B of the pre-registration: dynamical EWSB is not observed at these lattice parameters. The Higgs VEV $v = 246 \text{ GeV}$ remains [IMPOSED] in the parametric-insertions catalogue.

6.2 Three-Generation Cold-Start

Radial-flux seed with amplitude 0.2 across $L = 16$, evolved for 1000 ticks. Count of manifested species: zero of any sign. The radial configuration does not cross the genesis threshold; this is consistent with the Moore-layer-decomposition claim being topological (via $C(3, 2) = 3$) rather than a dynamical consequence of evolution.

6.3 Large- L $\alpha_{\text{large}L}$ Extrapolation Scan

α_{eff} measured at $L \in \{32, 48, 64\}$ via the tail-plateau method (§4). Fit $\alpha(L) = \alpha_{\text{large}L} + b/L^2$:

$L = 32$	$\alpha = 0.0352$
$L = 48$	$\alpha = 0.0140$ (outlier — non-power-of-2 lattice)
$L = 64$	$\alpha = 0.0331$
$\alpha_{\text{large}L}$	$= 0.0214, b = +10.9$

Finding. The $1/L^2$ fit value $\alpha_{\text{large}L} = 0.0214$ is a ratio of 2.94 to $\alpha_{\text{ref}} = 0.00730$, improving on the Phase-2C single-scale $L = 64$ value (0.033, ratio 4.5) but not yet at the pre-registered 1% precision. The $L = 48$ outlier destabilises the fit; multi-seed statistics and power-of-2-only scans are needed to give an empirical residual band on $\alpha_{\text{large}L}$.

7 Catalogue of Upgrades and Non-Upgrades

Of the ~ 162 entries in CATALOG_PARAMETRIC_INSERTIONS.md (49 cross-referenced across lepton, quark, meson, baryon, mixing, decay, precision-QED, neutrino, Higgs, QCD, cosmological, cross-section, and null-prediction sectors), this campaign upgrades exactly *zero* from [IMPOSED] or [PARAMETRIC] to [DERIVED]. Two new rows are added:

- “ α_{EM} running under blocking” — [MEASURED], scale-dependent with $\beta_{\text{meas}}/\beta_{\text{QED}} \approx -160$ (sign matches, magnitude off). Cross-references DERIV_BETA_FUNCTION_MEASURED.md.
- “Operator scaling dimensions” — [MEASURED], $\Delta \in [0.39, 0.75]$ for five operators, dominated by pulse-envelope artefact at the measurement scale. Cross-references DERIV_OPERATOR_SPECTRUM.md.

Neither row justifies reclassification of existing entries; both document new measurement output rather than derivation. The net effect on the catalog’s distribution of tags is: 23 [DERIVED]/[THEOREM], 131 [PARAMETRIC], 10 [IMPOSED]/[SELECTION], 2 [MEASURED] (Phase-2 and Phase-3 rows above).

8 Concrete Follow-Up Tickets

The five concrete extensions that would close quantitative gaps:

1. **$L = 128$ large- L extrapolation scan** (Phase 2/4 extension). At $L = 128$ the fit window $r \in [4, 32]$ extends $3\times$ past the Phase-2 screening length $\lambda = 10.6$, entering the pure- Coulomb regime. Expected runtime: ~ 10 minutes with existing infrastructure.
2. **`gauss_project_converged()` extension**. Replace the fixed 6-iteration SOR with an iterate-to-tolerance version, gated by a new toggle to avoid hot-path regression. Phase-1C Ward-identity measurement would then reach $\leq 10^{-6}$.
3. **Multi-seed statistics**. Eight independent seeds at $L = 48$ would resolve whether the Phase-4C outlier is statistical or systematic.
4. **Higher-amplitude cold-start**. EWSB with initial $\langle |\mathbf{J}| \rangle \geq 0.5$ (genesis-threshold-crossing) to test whether Branch A of Phase-4A is reachable.
5. **Confinement-era operator scan**. Re-run Phase 3’s six- operator spectrum on **flux-baryon** or **qcd-confinement** scenarios where long-range structure exists without a spatial envelope contaminating the decay.

Each ticket is self-contained, fits existing module boundaries, and requires no new physics assumptions.

9 Post-Campaign Follow-Up Measurements

After the first draft of this manuscript was committed, the five follow-up tickets enumerated in §8 were executed. The full narrative is in `DERIV_GAP_CLOSURE.md`; the key outcomes requiring manuscript integration:

T1 *gauss_project_converged* — stencil-mismatch floor. The helper `ftd::eft::gauss_project_converged` was implemented and applied to the Phase-1C charge-pair configuration. After 500 iteration cycles (~ 3000 SOR sweeps) the residual did not converge below the single-cycle value. Root cause: the engine’s SOR uses an 18-point Laplacian on φ , while the divergence in the Poisson source uses 6-point central difference. The two discretisations are not consistent, so even a perfectly-solved 18-pt Poisson equation does not cancel the 6-pt divergence, and repeated projection is not a contraction. The manuscript’s §8 recommendation to “ship a `gauss_project_converged()` extension” is accordingly superseded: the true fix requires either matched-stencil Poisson or a multigrid/conjugate-gradient solver, both engine-level changes outside this programme’s scope.

T2 multi-seed β ensemble — measurement is seed-robust. Four seeds (initial flux amplitudes $0.03, 0.05, 0.07, 0.10$) $\times$ three lattice sizes give α_{fit} variation $\leq 2.2\%$. The Phase-4C $L = 48$ outlier is therefore NOT statistical noise; it is a real systematic attributable to the non-power-of-2 periodic-image structure of $L = 48$. Future scans should use power-of-2 lattices only.

T3 $L = 128$ scan — “Yukawa screening” is finite-size. The three-scale Yukawa-length measurement gives $\lambda(L = 32) = 2.88$, $\lambda(L = 64) = 10.57$, $\lambda(L = 128) = 25.61$. This is $\lambda \approx L/5$: the screening length is a *fixed fraction of the lattice size*, not a physical Yukawa mass. As L grows through the tested range, λ grows in lock-step and pure Coulomb behaviour is approached at the largest tested L . This recontextualises the §4 β -measurement: the “screening” is at least partially a finite-size periodic-image cutoff, not physical running.

Using the slope-method comparison (which does not fit through the screening envelope), the $\alpha(L = 128)/\alpha(L = 64) = 1.09$ ratio gives $\Delta\alpha/\alpha = -0.08$ per blocking factor, or $\beta \approx -4 \times 10^{-3}$ at $g = 0.36$ — still negative, now about $80\times$ the one-loop QED prediction rather than $160\times$. The headline “new quantitative prediction” survives at reduced magnitude; a clean separation of physical running from finite-size screening still needs $L \geq 256$.

T4 EWSB amplitude sweep — Branch A at $\text{amp} \geq 0.80$! Sweeping the EWSB cold-start amplitude across $\{0.15, 0.30, 0.50, 0.80\}$:

amp	$\langle J \rangle_0$	$\langle J \rangle_f$	ratio	$ \Sigma s _f$
0.15	0.18	0.09	0.48	0
0.30	0.36	0.18	0.48	0
0.50	0.61	0.29	0.48	0
0.80	0.97	2.99	3.08	62

At $\text{amp} = 0.80$, $\langle |J| \rangle$ tripled and 62 charges spontaneously emerged from the vacuum. The genesis threshold is crossed somewhere in $(0.50, 0.80)$. This is the first dynamical manifestation event in the EFT programme, and it reopens the question of dynamical EWSB that Phase 4A closed pessimistically. Higgs-VEV reclassification to [DERIVED] is still not warranted (no principle selects $\text{amp} = 0.80$ as unique cold-start condition), but **Branch A now exists** for this engine.

T5 confinement operator scan — pulse-envelope artefact confirmed. Rerunning the six-operator basis on the **flux-baryon** scenario (three bound quarks with long-range confinement strings) instead of the smooth Gaussian pulse:

Operator	naive Δ	pulse Δ	flux-baryon Δ
JJ	2	0.531	0.448
$(\nabla \cdot J)^2$	4	0.458	1.690
$(\nabla \times J)^2$	4	0.391	0.676
$J \cdot J^4$	4	0.753	0.836

The $(\nabla \cdot J)^2$ dimension **jumps from 0.46 to 1.69** when we swap scenario — a $3.7\times$ shift moving this operator into the marginal bracket as the naive-counting pre-registration expected. The pulse-envelope-artefact hypothesis of §5 is [CONFIRMED]. The operator basis is physical; scenario choice mattered.

10 Day-2 Campaign: Four More Threads

After the gap-closure tickets (§9) and in response to the user’s direction to “close the gaps,” four additional threads ran. These are documented in full in `DERIV_DAY2_CAMPAIN.md`; the headline results:

Thread 2 — Matched-stencil CG Poisson (§9 T1 resolved). A Yee-style staggered-difference Poisson solver (backward divergence, forward gradient, composed to the standard 7-point nearest-neighbor Laplacian) was implemented in `engine/include/ftd/eft/matched_poisson.h`. On the canonical charge-pair test, the deep-vacuum Ward residual drops from $\sim 10^{-2}$ (engine’s SOR) to $\leq 10^{-8}$ — a million-fold improvement. The structural caveat documented in T1 (vacuum voxels adjacent to particles retain an $O(|\mathbf{J}|_p)$ boundary-layer residual because the engine convention leaves particle flux alone) is preserved. The Ward-identity precision floor on dynamical α -measurements is now software-limited only by the CG tolerance.

Thread 1b — EWSB amplitude threshold map. On $L = 32$ with 5000 ticks per amplitude, a clean first-order condensation transition is located between $\text{amp} = 0.6$ and $\text{amp} = 0.7$:

amp	$\langle J \rangle_f$	N^+ / N^-	charge frac	imbalance
0.50	0.33	0 / 0	0%	0
0.60	0.40	0 / 0	0%	0
0.70	2.26	15 776 / 16 992	100%	−1 216
0.80	2.26	17 478 / 15 290	100%	+2 188
0.90	2.70	20 167 / 12 601	100%	+7 566

At $L=32$ the condensation is lattice-wide: all 32 768 voxels manifest state simultaneously (contrast to the gap-closure $L=16$ result where only 62 of 4096 voxels manifested, 1.5%). The charge imbalance $N^+ - N^-$ grows monotonically with amplitude, consistent with a Higgs-like vacuum selecting a sign above threshold.

Thread 3 — Condensate spectroscopy. From the binary lattice dumps, the flux-flux and charge-charge correlators of the condensate both fit clean exponential decays. At $\text{amp} = 0.80$, **both channels give** $m \approx 0.18$: $m_{\text{flux}} = 0.181$ ($R^2 = 0.99$) and $m_{\text{charge}} = 0.186$ ($R^2 = 0.96$). The ratio 0.97 indicates a single quasi-particle species dominates both channels. This is not a Standard-Model identification — the ratio is 10% off $M_W/M_Z = 0.881$ — but it is a real mass-gap observation on a spontaneously condensed vacuum state.

Thread 4 — Rutherford scattering cross-validation. A +1 projectile scattered off a locked +1 target at impact parameters $b \in \{3, 4, 5, 6, 7, 8\}$ on $L = 32$ with $v_0 = 0.3$ gives a deflection angle $\theta(b)$ from which α is fitted. Results:

b	θ (deg)	$\alpha_{\text{extracted}}$	ratio to α_{ref}
3	1.18	0.0349	4.78
4	0.97	0.0385	5.27
5	0.83	0.0411	5.64
6	0.74	0.0438	6.00
7	0.67	0.0463	6.35
8	0.62	0.0488	6.68
mean		0.0422 ± 0.0047	5.79

Critical comparison: the static $V(r)$ tail-plateau method at the same $L = 32$ gave $\alpha = 0.035$. The Rutherford scattering method at the smallest impact parameter $b = 3$ gives $\alpha = 0.035$ —

identical to 3 decimals. Two completely independent methods (static fit vs dynamic scattering) converge on the same answer in the Coulomb-dominated regime. **This proves the $5\times$ gap between measured α_{eff} and α_{ref} at $L = 32$ is genuine engine physics, not a $V(r)$ -fit artefact.**

The α -versus- b growth at large b is consistent with the same periodic-image finite-size screening that Yukawa-fitted the $V(r)$ method: both probes see the screening envelope at comparable distances.

Thread 1a — $L = 256$ Coulomb-tail measurement. Thread 1a (originally the GPU-accelerated $L \geq 256$ scan) completed on CPU in ‘fast-big’ mode (reduced ticks, coarser r -step) after the CUDA build proved blocked by interacting CMake 4 + NVCC 13 + MSBuild escape bugs (documented in STATUS_CUDA_BUILD.md). At $L = 256$ the $V(r)$ probe extends to $r = 84$ — a physical distance we had never before measured:

L	r_{max}	$\alpha_r(r_{\text{max}})/\alpha_{\text{ref}}$
64	20	4.1
128	40	3.8
256	84	1.4

Note (post-Phase-F retraction). The $\alpha_r = 0.010$ value at $L = 256, r = 84$ quoted above was obtained at ticks=100 (the Day-2 fast-big CPU scan) and turned out to be under-equilibrated. The full-precision ticks=300 GPU re-run gave $\alpha_r(r = 82) = 0.0271$ (see §11). The table above is kept for historical transparency; the three-point $1/L$ fit that produced the $1.23 \times \alpha_{\text{ref}}$ headline is **retracted** in favor of the four-point Phase-F extrapolation and its audit (AUDIT_ALPHA_EXTRACTION.md).

Historical (superseded) 3-point fit, shown for transparency only:

scaling law	α_{largeL}	$\alpha_{\text{largeL}}/\alpha_{\text{ref}}$	R^2
$1/L$ (Coulomb-in-periodic-box)	0.0090	1.23	0.66
$1/L^2$ (standard lattice-artefact)	0.0157	2.15	0.52

The correct result is the Phase-F 4-point plateau (§11): the $1/L^2$ fit at the largest tested L gives $\alpha_{\text{largeL}} \in [3.35, 3.74] \times \alpha_{\text{ref}}$ in the engine’s accumulator convention (residual band from three scaling-law choices), or $\approx 1.8 \times \alpha_{\text{ref}}$ after the classical $\frac{1}{2}$ convention correction derived in AUDIT_ALPHA_EXTRACTION.md. The “ $5\times$ gap” narrative has not been resolved by Phase-F; a residual factor of ~ 1.8 remains open.

Catalog impact. Four new [MEASURED] rows added to CATALOG_PARAMETRIC_INSERTIONS.md: Ward floor $\leq 10^{-8}$, EWSB threshold $\in (0.6, 0.7)$, condensate mass gap ≈ 0.18 , Rutherford $\alpha = 0.042 \pm 0.005$. Zero rows upgraded from [IMPOSED] to [DERIVED]. The honest-reporting discipline holds.

11 Phase F: GPU unblock + 4-point large-L plateau

The Day-2 manuscript closed with a 3-point $1/L$ extrapolation ($L = 64, 128, 256$) giving $\alpha_{\text{largeL}} = 1.23\alpha_{\text{ref}}$ via a $1/L$ scaling law. Phase F tests that extrapolation by (i) moving the GPU build to WSL2 + Ubuntu 22.04 + CUDA 13.0 (sidesteps the Windows CMake 4 escape bug entirely), and (ii) re-running the sweep under internally-consistent conditions at four lattice sizes.

WSL2 unblocks CUDA at 30× speedup. Windows CMake 4.2 + NVCC 13.0 + MSBuild interact pathologically on the `-D"CMAKE_INTDIR=\"Release\""` flag. Linux builds clean. Four surgical include-hygiene fixes (`<stdint>`, `<stdio>`) and one `pimpl` guard made the engine compile on Ubuntu 22.04. Single-seed $L = 256$ full-precision β scan drops from **> 2h (CPU, never completed on Windows)** to **4m 40s** on the RTX 5090. Details in `STATUS_CUDA_BUILD.md`.

The retraction. The Day-2 3-point series used a **fast-big** CPU run with `ticks=100` that was under-equilibrated — the flux field had not reached steady-state Coulomb tail amplitude around the locked charges. The full-precision GPU re-run at $L = 256$ with `ticks=300` gave $\alpha_r(r = 82) = 0.0271$, not the 0.010 the fast-big measurement had reported. The $1.23 \times \alpha_{\text{ref}}$ headline was built on that bad number.

The corrected 4-point series. At internally-consistent conditions (`ticks=150`, GPU, same $r/L \approx 0.31$ slice in each lattice):

L	r_{max}	$\alpha_r(r_{\text{max}})$	ratio to α_{ref}
64	20	0.02959	4.05
128	40	0.02970	4.07
256	82	0.02717	3.72
384	124	0.02632	3.61

Three candidate scaling laws:

fit form	α_{largeL}	$\alpha_{\text{largeL}}/\alpha_{\text{ref}}$
$1/L$	0.02640	3.62
$1/L^2$	0.02730	3.74
free $1/L^p$ (best $p = 0.5$)	0.02446	3.35

The plateau is the result — with a convention caveat. The $1/L^2$ fit at the largest tested L gives $\alpha_{\text{largeL}} \approx 3.74 \times \alpha_{\text{ref}}$, with an empirical residual band $[3.35, 3.74] \times \alpha_{\text{ref}}$ across the three scaling laws in the engine’s native accumulator convention. The 4-point series is flat enough that the fit is insensitive to scaling-law choice. **FTD does not match textbook QED at the Coulomb tail at the largest tested L .**

A post-audit (`AUDIT_ALPHA_EXTRACTION.md`) verifies that the three codepaths feeding this number (`coupling_measurement.h`, `measure_v_of_r.h`, `benchmark_emergent_alpha.cpp`) are mutually consistent, and isolates a **factor-of-two energy convention**: the engine’s `field_energy` is $\sum |J|^2$ with no $\frac{1}{2}$, whereas classical electromagnetism carries the $\frac{1}{2}$. The cross-term in $V(r) = E_{\text{pair}} - E_{\text{self}}^+ - E_{\text{self}}^-$ picks up this factor, so the engine’s measured $V(r)$ is $2 \times$ the classical-convention interaction energy. Under the classical convention the plateau is $\alpha_{\text{largeL}}^{\text{classical}} \approx 0.013 \approx 1.8 \times \alpha_{\text{ref}}$.

Phase G resolution — the residual is the Green’s function. The three interpretations left open by the audit (real physics, kinetic normalization, lattice Green’s function) collapse to the third on explicit calculation. The engine’s Gauss law $\nabla \cdot J = s$ has no coupling constant, so

$$V_{\text{engine}}(r, L) = -2G_L(r), \quad \alpha_r(r, L) = 2rG_L(r), \quad (1)$$

where $G_L(r)$ is the periodic lattice Poisson Green’s function on the L^3 torus with the 7-point Laplacian. This is a zero-free-parameter prediction. Computing G_L by FFT and comparing against the measured α_r at $L \in \{32, 64, 128, 256, 384\}$ gives $R^2 = 1.0000$, median 0.07% relative error, maximum 0.43% error across 16 points at $L = 384$ in the Coulomb-tail regime. The “plateau at $3.6 \times \alpha_{\text{ref}}$ ” is simply the value of the parameter-free lattice kernel $2rG_L(r)|_{r/L \approx 0.31}$, and has no QED content to deviate from. Full proof and data fit in `DERIV_EMERGENT_COULOMB_GEOMETRIC.md`.

Where α actually lives (and does not) in FTD. Two places carry fine-structure content:

1. **Explicit coupling toggle** (`coulomb`, off-path for Phase F): force $F = -\alpha \nabla \phi$ uses hardcoded $\alpha = 1/137$ [`PARAMETRIC`].
2. **Master quadratic** $x^2 - 16G^{*2}x + 16G^{*3} = 0$ has root x_+ with $1/x_+ \approx \alpha_{\text{ref}}$ to sub-ppm; a pure number-theoretic identity independent of simulation [`THEOREM`].

The Phase-F $V(r)$ measurement — the `emergent_forces` code path — carries *neither*. Reproducing α as an emergent outcome of lattice dynamics requires Phase H: adding a single coupling $g_c = \sqrt{4\pi\alpha_{\text{ref}}}$ to the Gauss source and verifying that the measured α_r plateau now sits at α_{ref} . That measurement is spec’d but not yet run.

Modular pipeline. Phase F also shipped the `engine/sim/` subsystem: a `Pipeline<Backend>` template (CPU + GPU specialisations), 3+ pre-built `Observable<T>` classes, and 2 pipeline-based benchmark rewrites. Every future EFT measurement can be a ≤ 50 -LOC declarative program. GPU/CPU parity verified at permille after 1 tick, 1% after 500 ticks. Details in `DERIV_DAY2_CAMPAGN.md` and `~/claude/plans/vivid-marinating-pudding.md`.

12 Conclusion

The EFT Recovery Program closes with a Branch-A native effective field theory and a structurally clarified relationship to QED matching.

Branch A (native source/flux EFT) — complete at the minimum-viable level. The five pre-registered measurement pillars produced finite, reproducible signals on the native FTD source/flux variables. Phase 1 (field dictionary, action/measure, operator basis, blocking, Ward identities) and Phase 2 (β -function determination via flux-energy density across $b \in \{1, 2, 4, 8\}$) close at the level required by the bridge contract `SPEC_FTD_EFT_BRIDGE_CONTRACT.md`: the five gates 1–5 plus the Phase-2 measurement at $b = 8$ all PASS, with $|\beta_{\mathcal{E}}| \leq 0.08$ per b -decade — consistent with a Gaussian fixed point within 1σ . Together these results establish that FTD defines a native source/flux EFT with a Coulomb-like long-distance response, transverse propagating modes, and measurable Wilsonian flow of four coupling coefficients (C_L, K_T, Z_j, g_{sJ}).

Branch B (QED matching) — structurally decoupled. The original framing of $\alpha_{\text{eff}}(L)$ as a Wilsonian running has been retracted. Phase F + the post-audit `AUDIT_ALPHA_EXTRACTION.md` re-identifies the plateau as a zero-free-parameter periodic lattice Poisson Green’s function $\alpha_r(r, L) = 2rG_L(r)$ at $r/L \approx 0.31$ ($R^2 = 1.0000$ at $L = 384$, median 0.07% residual; full reframe in `DERIV_EMERGENT_COULOMB_G`). The four pre-registered routes to derive $1/\alpha$ from the master quadratic (R1 transverse stiffness, R2 source-current normalisation, R3 two-sector response eigenvalue, R4 projected Dirac) all closed

negative under the current projected action. The lead-physicist diagnosis is that Phase J ultralocality structurally decouples the algebraic spine (master quadratic, G_* , CM-curve uniqueness) from the dynamical EFT: the action data does not contain the polynomial data. Closure of the QED-matching branch likely requires a non-action injection mechanism (boundary conditions, observable selection, quantization choice) and is not session-tractable.

What is upgraded vs not upgraded. No claim in the FTD parametric-insertions catalogue is upgraded from [IMPOSED] to [DERIVED] by this campaign. Two new rows tag *measured* native-EFT phenomena (the b -decade $\beta_{\mathcal{E}}$ flow and operator scaling-dimension fits) that were previously unmeasured. The lattice-Coulomb-geometry plateau is correctly tagged [DERIVED] – it is a closed-form lattice Green’s function, not an empirical match. The catalogue is honest; the campaign did not retrofit tags.

The empirical match $x_+ \approx 1/\alpha$. The 1.26 ppm match between the master-quadratic root $x_+ \approx 137.036$ and CODATA α^{-1} remains [STRONGLY MOTIVATED CONJECTURE]. **Correction (2026-08-05):** an earlier version of this paragraph called the supporting structural evidence “substantial,” citing the uniqueness scans and a $\sim 4 \times 10^5 : 1$ figure. The scans are withdrawn as evidence: the adversarial scan sits exactly at its chance base rate (null expects 1.42–1.67 matchers, 1 found; FTD-0791); the extended scan failed its own pre-registered base-rate control ($N_{\text{null}} = 0.0014$ — a zero count is the chance expectation) and its published count was wrong (four cubic dual-matchers exist in the declared family; FTD-0802); and the Chowla–Selberg $h \geq 2$ null is a range artifact — the full per-ideal-class Damerell scan finds 1,271 fundamental dual-matchers by $|d| \leq 500,000$ (FTD-0321). The 4×10^5 figure was retracted 2026-06-24. What stands: the Eisenstein-family null, and the BCC complex-structure theorem (DERIV_BCC_COMPLEX_STRUCTURE.md, FTD-0122), which unifies two of the four order-four occurrences in cubic-lattice algebra via a natural $\mathbb{Z}[i]$ -module structure. The identification presently has no numerical-uniqueness support; its support is structural. A formal axioms-to- α derivation is structurally blocked, not merely unsolved.

Net characterisation. FTD is, on the present evidence, a native Wilsonian EFT with measured long-distance Coulomb-like response, transverse wave modes, native RG flow, and a structurally decoupled but empirically suggestive relationship to QED. The first publishable claim under the bridge contract is the native-EFT result, not an α -derivation. That claim is what this paper supports.

A Reproducibility

```
# Build
cmake -S engine -B engine/build -DCMAKE_BUILD_TYPE=Release
cmake --build engine/build --config Release --target \
    test_eft_anisotropy benchmark_lorentz_recovery \
    test_eft_ward_identity test_eft_blocking \
    benchmark_beta_function test_eft_operator_spectrum \
    benchmark_dynamical_sm

# All six Phase 1-4 tests
cd engine/build && ctest -C Release -L eft --output-on-failure
```

```
# Full beta measurement + analyzer
python scripts/benchmarks/measure_beta_function.py
```

```
# Full dynamical-SM campaign (~80s)
./engine/build/Release/benchmark_dynamical_sm.exe
```

Output files: `scripts/benchmarks/results/eft_beta/{beta_raw.csv, beta_results.json, beta_report.md}`.

B Module inventory

EFT-specific code shipped in this campaign (all under new `eft/` namespaces, no regression to existing modules):

- `engine/include/ftd/eft/anisotropy.h` (263 LOC)
- `engine/include/ftd/eft/lorentz_recovery.h` (210 LOC)
- `engine/include/ftd/eft/ward_identities.h` (209 LOC)
- `engine/include/ftd/eft/blocking.h` + `src/eft/blocking.cpp` (145 + 270 LOC)
- `engine/include/ftd/eft/coupling_measurement.h` (195 LOC)
- `engine/include/ftd/eft/operator_spectrum.h` (245 LOC)
- `engine/src/eft/qcd_one_loop_perturbative.cpp` (52 LOC — renamed from `ontic_running_coupling.cpp` [IMPOSED]-tagged)
- Six test executables, ~1200 LOC
- One Python analyzer, 461 LOC
- Five theory documents, ~1900 LOC combined

References

- [1] FTD specification. `docs/SPEC_FTD.md`, FTD project repository, 2026.
- [2] Complete Standard Model computation from FTD integers. `scripts/proofs/proof_complete_sm.py`, 2026.
- [3] Master quadratic and the empirical identification $1/x_+ \approx \alpha$. `docs/theory/03_derivations/DERIV_MASTER_QUADRATIC_GAP_EQUATION.md`, 2026.
- [4] Specification of FTD’s nine algebraic-spine theorems. `docs/theory/01_reference/SPEC_ALGEBRAIC_SPINE.md`, 2026.
- [5] W. J. Steinmetz III and C. Pacitanimulli, *A π -free generator of the lemniscatic field, with applications to a closed-form quadratic for the fine-structure constant*, companion paper / preprint, 2026.

- [6] W. J. Steinmetz III and C. Pacitanimulli, *A $\mathbb{Z}[i]$ -module structure on body-centered-cubic corners*, companion paper / preprint, 2026.
- [7] $N_c = 3$ from topology: four independent routes. docs/theory/03_derivations/DERIV_NC_FROM_TOPOLOGY.md, 2026.
- [8] Moore Layer Decomposition Theorem. docs/theory/08_structural/THEOREM_MOORE_LAYER_DECOMPOSITION.md, 2026.
- [9] Continuum limit: FTD = lattice QED conditionally. docs/theory/03_derivations/DERIV_CONTINUUM_LIMIT_QED_EQUIVALENCE.md, 2026.
- [10] EFT Recovery Program pre-registration. docs/theory/10_eft_program/SPEC_EFT_RECOVERY_PROGRAM.md, 2026.
- [11] Phase 1 symmetry-recovery measurements. docs/theory/10_eft_program/DERIV_SYMMETRY_RECOVERY.md, 2026.
- [12] Phase 2 measured β -function. docs/theory/10_eft_program/DERIV_BETA_FUNCTION_MEASURED.md, 2026.
- [13] Phase 3 operator spectrum. docs/theory/10_eft_program/DERIV_OPERATOR_SPECTRUM.md, 2026.
- [14] Phase 4 dynamical SM emergence tests. docs/theory/10_eft_program/DERIV_DYNAMICAL_SM_EMERGENCE.md, 2026.
- [15] Parametric insertions catalogue. docs/theory/07_assessment/CATALOG_PARAMETRIC_INSERTIONS.md, 2026.
- [16] L. P. Kadanoff, Scaling laws for Ising models near T_c , *Physics* **2**, 263 (1966); K. G. Wilson, Renormalization group and critical phenomena, *Phys. Rev. B* **4**, 3174 (1971).